

Method of Stages: Non-Exponential Distributions

EdgeBasedModels.jl Vignette 10

Simon Frost

2026-03-30

Table of contents

Introduction	1
Setup	2
Exponential vs Erlang distributions	2
Building an Erlang SIR model	3
Comparing exponential vs Erlang epidemics	5
Varying the number of stages	8
GammaApproxStage: matching a target CV	9
SEIR with Erlang stages	11
Summary	15

Introduction

Standard edge-based compartmental models assume exponentially distributed sojourn times in each disease compartment. The exponential distribution is memoryless, leading to simple ODE systems, but it is often unrealistic: real infectious periods tend to be concentrated around a mean value rather than broadly spread.

The method of stages (also called the linear chain trick) offers a simple fix. Replace a single compartment with a chain of n identical sub-compartments, each with rate $n\gamma$. The total sojourn time then follows an Erlang distribution $\text{Erlang}(n, n\gamma)$, which has:

- Mean: $1/\gamma$ (same as the original exponential)
- Variance: $1/(n\gamma^2)$ (decreasing with n)
- Coefficient of variation: $\text{CV} = 1/\sqrt{n}$

As $n \rightarrow \infty$, the Erlang distribution concentrates at the mean, approaching a fixed (deterministic) sojourn time. Because the mean is preserved, the basic reproduction number R_0 and the final epidemic size remain unchanged — only the transient dynamics (peak height and timing) are affected.

This approach was described by Lloyd (2001) for compartmental models and by Sherborne et al. (2018) for edge-based models on networks.

Setup

```
using EdgeBasedModels
using ModelingToolkit
using OrdinaryDiffEq
using Symbolics
using Plots
```

```
@parameters β γ κ
pgf = poisson_pgf(κ)
κ_val = 5.0
β_val = 0.5
γ_val = 0.1
ψ(x) = exp(κ_val * (x - 1))
tspan = (0.0, 120.0)
```

(0.0, 120.0)

Exponential vs Erlang distributions

Before building epidemic models, let us visualize how the Erlang distribution compares to the exponential. All distributions below have the same mean sojourn time $1/\gamma = 10$, but different shapes.

```
function erlang_pdf(x, n, total_rate)
    λ = n * total_rate
    return λ^n * x^(n - 1) * exp(-λ * x) / factorial(n - 1)
end

x = range(0, 40, length = 500)
p = plot(xlabel = "Sojourn time", ylabel = "Density",
```

```

    title = "Erlang distributions (mean =  $1/\gamma_{\text{val}}$ )"
for (n, ls) in [(1, :solid), (2, :dash), (5, :dashdot), (20, :dot)]
    label = n == 1 ? "Exp( $\gamma$ ) [CV=1.00]" : "Erlang( $n$ ,  $n\gamma$ ) [CV= $\text{round}(1/\sqrt{n}; \text{digits}=2)$ ]"
    plot!(p, x, erlang_pdf.(x, n,  $\gamma_{\text{val}}$ ), label = label, linewidth = 2, linestyle = ls)
end
vline!(p, [1 /  $\gamma_{\text{val}}$ ], label = "Mean =  $1/\gamma$ ", color = :black, linestyle = :dash, linewidth = 2)
p

```




Figure 1: Erlang distributions with the same mean ($1/\gamma = 10$) but increasing shape parameter n . As n increases, the distribution concentrates around the mean.

The exponential ($n = 1$) has $CV = 1$ — its standard deviation equals its mean. As n increases, the distribution narrows: by $n = 20$, the CV is only 0.22, and the density is tightly concentrated around the mean of 10.

Building an Erlang SIR model

EdgeBasedModels provides `ErlangStage` to define a stage with Erlang-distributed sojourn time, and `expand_erlang_stages` to convert it into a standard `DiseaseProgression`.

An `ErlangStage(:I, n,  $\gamma$ ; transmission_rate =  $\beta$ )` represents an infectious period that is Erlang-distributed with shape n and overall rate γ . When expanded, it produces n sub-stages $I_1, I_2, \dots, I_n$, each inheriting the transmission rate β and connected by transitions at rate $n\gamma$:

$$I_1 \xrightarrow{n\gamma} I_2 \xrightarrow{n\gamma} \dots \xrightarrow{n\gamma} I_n \xrightarrow{n\gamma} R$$

```
# Standard exponential SIR (n = 1)
model_exp = build_sir(pgf,  $\beta$ ,  $\gamma$ ; form = :expanded)
println("Exponential SIR: ", length(ModelingToolkit.equations(model_exp.system)), " ODEs")
```

Exponential SIR: 4 ODEs

```
# Erlang SIR with n = 5 sub-stages for the infectious period
n_stages = 5
erlang_I = ErlangStage(:I, n_stages,  $\gamma$ ; transmission_rate =  $\beta$ )
prog_erlang = expand_erlang_stages(
    [erlang_I, DiseaseStage(:R)],
    [DiseaseTransition(:I, :R, n_stages *  $\gamma$ )];
    entry = :I,
)
model_erlang = build_edge_system(
    StaticConfigurationModel(pgf, prog_erlang);
    form = :expanded,
)
println("Erlang SIR (n=$n_stages): ", length(ModelingToolkit.equations(model_erlang.system))
```

Erlang SIR (n=5): 8 ODEs

i Note

The exit transition from I to R must use the sub-stage rate $n\gamma$, not the overall rate γ . This ensures all sub-stages — including the last one — have identical sojourn-time distributions, which is required for the total to be Erlang.

Let us inspect the expanded ODE system:

```
for eq in ModelingToolkit.equations(model_erlang.system)
    println(" ", eq)
end
```

```

Differential(t, 1)(R(t)) ~ 5(1 - exp((-1 + θ(t))*κ) - R(t))*γ
Differential(t, 1)(phi_R(t)) ~ 5phi_I_5(t)*γ
Differential(t, 1)(phi_I_5(t)) ~ 5phi_I_4(t)*γ - phi_I_5(t)*β - 5phi_I_5(t)*γ
Differential(t, 1)(phi_I_4(t)) ~ 5phi_I_3(t)*γ - phi_I_4(t)*β - 5phi_I_4(t)*γ
Differential(t, 1)(phi_I_3(t)) ~ -phi_I_3(t)*β - 5phi_I_3(t)*γ + 5phi_I_2(t)*γ
Differential(t, 1)(phi_I_2(t)) ~ -phi_I_2(t)*β - 5phi_I_2(t)*γ + 5phi_I_1(t)*γ
Differential(t, 1)(phi_I_1(t)) ~ -phi_I_1(t)*β - 5phi_I_1(t)*γ + phi_S(t)*(phi_I_3(t) + phi_I_2(t) + phi_I_1(t))
Differential(t, 1)(θ(t)) ~ -(phi_I_3(t) + phi_I_2(t) + phi_I_4(t) + phi_I_5(t) + phi_I_1(t))*γ

```

Comparing exponential vs Erlang epidemics

Both models have the same mean infectious period ($1/\gamma$) and per-edge transmission rate (β), so R_0 is identical. Let us compare the dynamics.

```

# Solve exponential SIR
ic_exp = default_initial_conditions(model_exp)
prob_exp = ODEProblem(
    model_exp.system,
    merge(ic_exp, Dict{β => β_val, γ => γ_val, κ => κ_val}),
    tspan,
)
sol_exp = solve(prob_exp, Tsit5(); saveat = 0.5)

state_names_exp = string.(ModelingToolkit.unknowns(model_exp.system))
θ_idx_exp = findfirst(s → startswith(s, "θ"), state_names_exp)
R_idx_exp = findfirst(s → startswith(s, "R"), state_names_exp)
θ_exp = sol_exp[θ_idx_exp, :]
R_exp = sol_exp[R_idx_exp, :]
S_exp = ψ.(θ_exp)
I_exp = 1.0 .- S_exp .- R_exp

```

```

241-element Vector{Float64}:
 0.00498752080731768
 0.0047442765291545504
 0.0045128954494401685
 0.004292798828178263
 0.004083436019791223
 0.00388428450359447
 0.0036948461339441543
 0.0035146465844452203
 0.003343234715160879

```

```

0.0031801816282027466
?
6.211918420536561e-8
6.187977976811893e-8
6.153985682762902e-8
6.103506919926532e-8
6.029835240145553e-8
5.925992364961413e-8
5.784728186047916e-8
5.598520765124487e-8
5.359576333782701e-8

```

```

# Solve Erlang SIR
ic_erlang = default_initial_conditions(model_erlang)
prob_erlang = ODEProblem(
    model_erlang.system,
    merge(ic_erlang, Dict{ $\beta$   $\Rightarrow$   $\beta_{val}$ ,  $\gamma$   $\Rightarrow$   $\gamma_{val}$ ,  $\kappa$   $\Rightarrow$   $\kappa_{val}$ }),
    tspan,
)
sol_erlang = solve(prob_erlang, Tsit5(); saveat = 0.5)

state_names_erlang = string.(ModelingToolkit.unknowns(model_erlang.system))
 $\theta_{idx\_erl}$  = findfirst(s  $\rightarrow$  startswith(s, " $\theta$ "), state_names_erlang)
 $R_{idx\_erl}$  = findfirst(s  $\rightarrow$  startswith(s, "R"), state_names_erlang)
 $\theta_{erl}$  = sol_erlang[ $\theta_{idx\_erl}$ , :]
 $R_{erl}$  = sol_erlang[ $R_{idx\_erl}$ , :]
 $S_{erl}$  =  $\psi$ .( $\theta_{erl}$ )
 $I_{erl}$  = 1.0 .-  $S_{erl}$  .-  $R_{erl}$ 

```

```

241-element Vector{Float64}:
0.00498752080731768
0.0038842850659281617
0.003025084461361214
0.002355934190755476
0.0018347991295662963
0.001428942517067842
0.00111285684093625
0.0008666800893511618
0.000674998253958393
0.0005256350977320993
?
-3.864336028561871e-6

```

```

-3.6590918507913364e-6
-2.822148442652682e-6
-1.4574009076602154e-6
 2.5330181381826716e-7
 2.0501569445649515e-6
 3.595407870505933e-6
 4.473344140716422e-6
 4.1903014674207414e-6

```

```

plot(sol_exp.t, I_exp, label = "I (exponential)", linewidth = 2, color = :red)
plot!(sol_erlang.t, I_erl, label = "I (Erlang, n=5)", linewidth = 2, color = :blue)
plot!(sol_exp.t, R_exp, label = "R (exponential)", linewidth = 1, color = :red,
      linestyle = :dash)
plot!(sol_erlang.t, R_erl, label = "R (Erlang, n=5)", linewidth = 1, color = :blue,
      linestyle = :dash)
xlabel!("Time")
ylabel!("Fraction of population")
title!("Exponential vs Erlang SIR ( $\beta=\beta_{val}$ ,  $\gamma=\gamma_{val}$ )")

```




Figure 2: Exponential vs Erlang ($n=5$) SIR. The Erlang model produces a sharper, higher, earlier peak but converges to the same final size.

The final sizes converge to the same value — both epidemics infect the same total fraction of the

population. However, the Erlang model produces a sharper, higher peak that arrives earlier. This is because a more concentrated infectious period means individuals transmit infection over a narrower time window, leading to more synchronised dynamics.

Varying the number of stages

Let us sweep over $n = 1, 2, 5, 10, 20$ to see how the epidemic curve changes as the infectious period distribution narrows.

```
p = plot(xlabel = "Time", ylabel = "Fraction infected",
         title = "SIR with Erlang infectious period (varying n)")

for n in [1, 2, 5, 10, 20]
    if n == 1
        model_n = build_sir(pgf,  $\beta$ ,  $\gamma$ ; form = :expanded)
    else
        erl = ErlangStage(:I, n,  $\gamma$ ; transmission_rate =  $\beta$ )
        prog_n = expand_erlang_stages(
            [erl, DiseaseStage(:R)],
            [DiseaseTransition(:I, :R, n *  $\gamma$ )];
            entry = :I,
        )
        model_n = build_edge_system(
            StaticConfigurationModel(pgf, prog_n);
            form = :expanded,
        )
    end

    ic_n = default_initial_conditions(model_n)
    prob_n = ODEProblem(
        model_n.system,
        merge(ic_n, Dict( $\beta \Rightarrow \beta\_val$ ,  $\gamma \Rightarrow \gamma\_val$ ,  $\kappa \Rightarrow \kappa\_val$ )),
        tspan,
    )
    sol_n = solve(prob_n, Tsit5(); saveat = 0.5)

    sn = string(ModelingToolkit.unknowns(model_n.system))
     $\theta\_n$  = sol_n[findfirst(s  $\rightarrow$  startswith(s, " $\theta$ "), sn), :]
    R_n = sol_n[findfirst(s  $\rightarrow$  startswith(s, "R"), sn), :]
    I_n = 1.0 .-  $\psi$ .( $\theta\_n$ ) .- R_n

    cv = round(1 /  $\sqrt{n}$ ; digits = 2)
```



```

plot!(p, sol_n.t, I_n, label = "n=$n (CV=$cv)", linewidth = 2)
end
p

```




Figure 3: Effect of the number of Erlang stages on the infected curve. More stages $\Rightarrow$ sharper peak, same final size.

As n increases, the epidemic peak becomes taller and narrower. The curve approaches the limiting behaviour of a model with a fixed infectious period of exactly $1/\gamma$. All curves converge to the same final size, confirming that R_0 — and hence the final size — depends only on the mean infectious period, not its variance.

GammaApproxStage: matching a target CV

Instead of choosing n directly, it can be more natural to specify the desired coefficient of variation (CV) of the sojourn time. GammaApproxStage does this automatically by selecting:

$$n = \text{round}(1/\text{CV}^2)$$

This approximates a Gamma distribution with the specified mean and CV using an Erlang distribution with the closest integer shape parameter.

```
# Demonstrate the mapping from CV to number of sub-stages
println("CV → n sub-stages (mean sojourn =  $\frac{1}{\gamma_{\text{val}}}$ )")
println("-" ^ 35)
for cv in [1.0, 0.7, 0.5, 0.3, 0.2, 0.1]
    stage = GammaApproxStage(:I, 1 /  $\gamma_{\text{val}}$ , cv; transmission_rate =  $\beta_{\text{val}}$ )
    actual_cv = round(1 /  $\sqrt{\text{stage.n\_substages}}$ ; digits = 3)
    println("  CV = $cv → n = $(stage.n\_substages) (actual CV = $actual_cv)")
end
```

CV → n sub-stages (mean sojourn = 10.0)

```
CV = 1.0 → n = 1   (actual CV = 1.0)
CV = 0.7 → n = 2   (actual CV = 0.707)
CV = 0.5 → n = 4   (actual CV = 0.5)
CV = 0.3 → n = 11  (actual CV = 0.302)
CV = 0.2 → n = 25  (actual CV = 0.2)
CV = 0.1 → n = 100 (actual CV = 0.1)
```

Tip

Use `GammaApproxStage` when you have empirical estimates of the mean and CV of the infectious period (e.g., from clinical data), and `ErlangStage` when you want direct control over the number of sub-stages.

```
p = plot(xlabel = "Time", ylabel = "Fraction infected",
         title = "SIR parameterised by CV of infectious period")

for cv in [1.0, 0.5, 0.3, 0.2]
    # Use GammaApproxStage to determine n, then build with symbolic params
    n = GammaApproxStage(:I, 1 /  $\gamma_{\text{val}}$ , cv).n\_substages
    erl_cv = ErlangStage(:I, n,  $\gamma$ ; transmission_rate =  $\beta$ )
    prog_cv = expand_erlang_stages(
        [erl_cv, DiseaseStage(:R)],
        [DiseaseTransition(:I, :R, n *  $\gamma$ )];
        entry = :I,
    )
    model_cv = build_edge_system(
        StaticConfigurationModel(pgf, prog_cv);
        form = :expanded,
    )
    ic_cv = default_initial_conditions(model_cv)
```

```

prob_cv = ODEProblem(
    model_cv.system,
    merge(ic_cv, Dict{β ⇒ β_val, γ ⇒ γ_val, κ ⇒ κ_val}),
    tspan,
)
sol_cv = solve(prob_cv, Tsit5(); saveat = 0.5)

sn = string(ModelingToolkit.unknowns(model_cv.system))
θ_cv = sol_cv[findfirst(s → startswith(s, "θ"), sn), :]
R_cv = sol_cv[findfirst(s → startswith(s, "R"), sn), :]
I_cv = 1.0 .- ψ.(θ_cv) .- R_cv

plot!(p, sol_cv.t, I_cv, label = "CV=$cv (n=$n)", linewidth = 2)
end
p

```





Figure 4: Epidemic curves parameterised by the CV of the infectious period.

SEIR with Erlang stages

Both the latent (E) and infectious (I) periods can independently use Erlang distributions. Here we model an SEIR epidemic with Erlang-distributed E ($n_E = 3$ sub-stages) and Erlang-distributed I

($n_I = 5$ sub-stages):



```
@parameters σ
n_E = 3
n_I = 5

erlang_E = ErlangStage(:E, n_E, σ; transmission_rate = 0)
erlang_I_seir = ErlangStage(:I, n_I, γ; transmission_rate = β)

prog_seir_erlang = expand_erlang_stages(
  [erlang_E, erlang_I_seir, DiseaseStage(:R)],
  [
    DiseaseTransition(:E, :I, n_E * σ),
    DiseaseTransition(:I, :R, n_I * γ),
  ];
  entry = :E,
)

model_seir_erlang = build_edge_system(
  StaticConfigurationModel(pgf, prog_seir_erlang);
  form = :expanded,
)
println("Erlang SEIR: ", length(ModelingToolkit.equations(model_seir_erlang.system)), " ODEs")
```

Erlang SEIR: 11 ODEs

For comparison, the standard exponential SEIR:

```
model_seir_exp = build_seir(pgf, σ, β, γ)
println("Exponential SEIR: ", length(ModelingToolkit.equations(model_seir_exp.system)), " ODEs")
```

Exponential SEIR: 5 ODEs

```
σ_val = 0.2 # mean latent period = 5 days

# Solve exponential SEIR
ic_seir_exp = default_initial_conditions(model_seir_exp)
prob_seir_exp = ODEProblem(
```

```

    model_seir_exp.system,
    merge(ic_seir_exp, Dict( $\beta \Rightarrow \beta\_val$ ,  $\gamma \Rightarrow \gamma\_val$ ,  $\sigma \Rightarrow \sigma\_val$ ,  $\kappa \Rightarrow \kappa\_val$ )),
    tspan,
)
sol_seir_exp = solve(prob_seir_exp, Tsit5(); saveat = 0.5)

sn_se = string(ModelingToolkit.unknowns(model_seir_exp.system))
 $\theta\_se$  = sol_seir_exp[findfirst(s  $\rightarrow$  startswith(s, " $\theta$ "), sn_se), :]
R_se = sol_seir_exp[findfirst(s  $\rightarrow$  startswith(s, "R"), sn_se), :]
S_se =  $\psi$ .( $\theta\_se$ )
EI_se = 1.0 .- S_se .- R_se

```

241-element Vector{Float64}:

```

0.00498752080731768
0.0047442765271789675
0.0045128954484631376
0.004292798850391892
0.004083435971888061
0.0038842843802593743
0.003694846041310686
0.0035146465838806424
0.00334323494085032
0.0031801816899232407
␣
5.818145546593323e-8
5.6189005359121025e-8
5.427839099014786e-8
5.2420291235359684e-8
5.058429138841275e-8
4.8738883158538915e-8
4.685146467141299e-8
4.488834047088747e-8
4.281472151552307e-8

```

```

# Solve Erlang SEIR
ic_seir_erl = default_initial_conditions(model_seir_erlang)
prob_seir_erl = ODEProblem(
    model_seir_erlang.system,
    merge(ic_seir_erl, Dict( $\beta \Rightarrow \beta\_val$ ,  $\gamma \Rightarrow \gamma\_val$ ,  $\sigma \Rightarrow \sigma\_val$ ,  $\kappa \Rightarrow \kappa\_val$ )),
    tspan,
)

```

```

sol_seir_erl = solve(prob_seir_erl, Tsit5(); saveat = 0.5)

sn_sel = string.(ModelingToolkit.unknowns(model_seir_erlang.system))
θ_sel = sol_seir_erl[findfirst(s → startswith(s, "θ"), sn_sel), :]
R_sel = sol_seir_erl[findfirst(s → startswith(s, "R"), sn_sel), :]
S_sel = ψ.(θ_sel)
EI_sel = 1.0 .- S_sel .- R_sel

```

241-element Vector{Float64}:

```

0.00498752080731768
0.002131732043857496
0.0009111445611480057
0.0003894353537581231
0.00016646603946954795
7.119891080761137e-5
3.036189659632628e-5
1.295819578973293e-5
5.606061643502816e-6
2.3290472590887026e-6
␣
-3.94197407106462e-6
-3.343180877428756e-6
3.6093133108511255e-6
3.6599397101747325e-6
-3.3783189967266647e-6
-3.912234521440475e-6
2.738193365969463e-6
4.611537840814271e-6
1.9722442448520325e-6

```

```

plot(sol_seir_exp.t, EI_se, label = "E+I (exponential)", linewidth = 2, color = :red)
plot!(sol_seir_erl.t, EI_sel, label = "E+I (Erlang)", linewidth = 2, color = :blue)
plot!(sol_seir_exp.t, R_se, label = "R (exponential)", linewidth = 1, color = :red,
      linestyle = :dash)
plot!(sol_seir_erl.t, R_sel, label = "R (Erlang)", linewidth = 1, color = :blue,
      linestyle = :dash)
xlabel!("Time")
ylabel!("Fraction of population")
title!("Exponential vs Erlang SEIR (nE=$n_E, nI=$n_I)")

```




Figure 5: Exponential vs Erlang SEIR. Erlang stages produce sharper dynamics for both the latent and infectious periods.

The Erlang SEIR produces a sharper epidemic curve with a higher, earlier peak. As with SIR, the final epidemic size is the same for both models since it depends only on R_0 and the degree distribution.

Summary

- The method of stages replaces a single compartment with n identical sub-compartments, giving an Erlang-distributed sojourn time with the same mean but reduced variance.
- Same mean $\Rightarrow$ same R_0 and final epidemic size; different variance $\Rightarrow$ different peak timing and height.
- More stages $\Rightarrow$ sharper, more peaked epidemics, approaching the fixed-duration limit.
- `ErlangStage(name, n,  $\gamma$ )` creates an Erlang stage with n sub-stages and overall rate γ .
- `GammaApproxStage(name, mean, cv)` automatically selects $n = \text{round}(1/\text{CV}^2)$ to match a target coefficient of variation.
- `expand_erlang_stages` converts Erlang stages into a standard `DiseaseProgression` compatible with the full edge-based framework.
- Multiple compartments (E , I , etc.) can each independently use Erlang distributions.