

Table of contents

Clustering and Triangle Motifs	1
Introduction	1
Setup	2
Clustered networks	2
Building a clustered SIR model	3
Comparing clustered vs unclustered epidemics	4
Varying the clustering coefficient	8
Clustering coefficient and R_0	8
SEIR with clustering	10
Summary	13
Simulation validation	14

Clustering and Triangle Motifs

Simon Frost 2026-05-14

- [Introduction](#)
- [Setup](#)
- [Clustered networks](#)
- [Building a clustered SIR model](#)
- [Comparing clustered vs unclustered epidemics](#)
- [Varying the clustering coefficient](#)
- [Clustering coefficient and \$R_0\$](#)
- [SEIR with clustering](#)
- [Summary](#)
- [Simulation validation](#)

Introduction

Real contact networks exhibit clustering — the tendency for two people who share a mutual contact to also be connected to each other, forming triangles. This property, also called triadic closure, is ubiquitous in social networks but absent from the standard configuration model in the large- N limit.

Clustering has important epidemiological consequences. When transmission can travel around a triangle, one of the three edges is “redundant”: if A infects B and B infects C , the edge from A to C cannot cause a *new* infection. This correlation means that clustering reduces epidemic severity compared to a tree-like network with the same degree distribution.

The edge-based compartmental model framework handles clustering by introducing a bivariate probability generating function $g(x, y)$ where x tracks single-edge stubs and y tracks triangle-edge stubs. This approach follows Volz (2011) and Miller (2009).

In this vignette we show how to:

1. Build clustered network models using ClusteredPGF
2. Compare clustered and unclustered epidemics
3. Explore how increasing clustering reduces R_0 and epidemic severity

Setup

```
using EdgeBasedModels
using ModelingToolkit
using OrdinaryDiffEq
using Symbolics
using Plots
```

Clustered networks

In a clustered network each node has two types of half-edges:

- Single-edge stubs (s): connect to a random partner (tree-like)
- Triangle-edge stubs (t): connect in groups of three to form triangles

The joint degree distribution is encoded by a bivariate PGF:

$$g(x, y) = \sum_{s,t} p_{s,t} x^s y^t$$

where $p_{s,t}$ is the probability that a node has s single-edge stubs and t triangle-edge stubs. The mean single degree is $\langle s \rangle = g_x(1, 1)$ and the mean triangle degree is $\langle t \rangle = g_y(1, 1)$.

The clustering coefficient measures the fraction of connected triples that are closed into triangles:

$$C = \frac{2\langle t \rangle}{2\langle t \rangle + \langle s \rangle}$$

For a Poisson clustered network with mean single-edge degree κ_s and mean triangle-edge degree κ_t :

$$g(x, y) = e^{\kappa_s(x-1) + \kappa_t(y-1)}$$

```
κ_s = 3.0
κ_t = 1.0
cpgf = clustered_poisson_pgf(κ_s, κ_t)
println("Mean single-edge degree: ", mean_single_degree(cpgf))
println("Mean triangle-edge degree: ", mean_triangle_degree(cpgf))
println("Clustering coefficient: ", clustering_coefficient(cpgf))
```

```
Mean single-edge degree: 3.0exp(0.0)
Mean triangle-edge degree: exp(0.0)
Clustering coefficient: 0.4
```

Each triangle-edge stub participates in one triangle with two partners, so the total mean degree (number of distinct neighbours) is $\langle s \rangle + 2\langle t \rangle$.

Building a clustered SIR model

The key insight of the clustered EBCM is that we need two survival probabilities:

- $\theta_2(t)$: probability a single-edge partner has not transmitted infection
- $\theta_3(t)$: probability a triangle-edge partner has not transmitted infection

The susceptible fraction is then $S(t) = g(\theta_2, \theta_3^2)$, where the square on θ_3 accounts for the two triangle partners per triangle stub.

Let us build both an unclustered and a clustered SIR model with comparable mean degrees.

```
@parameters β γ

# Universal anchors: γ=0.25, R₀=2, β derived per scenario (see plan.md)
γ_val = 0.25
R0_target = 2.0
κ_total = 5.0
T_unclustered = R0_target / κ_total
β_val = T_unclustered * γ_val / (1 - T_unclustered)

# Unclustered: all edges are single-stubs, mean degree 5
pgf_unclustered = poisson_pgf(κ_total)
model_unclustered = build_sir(pgf_unclustered, β, γ; form = :expanded)
```

```
# Clustered:  $\kappa_s=3$ ,  $\kappa_t=1 \rightarrow$  total mean degree =  $3 + 2(1) = 5$ 
cp_gf = clustered_poisson_pgf(3.0, 1.0)
model_clustered = build_clustered_sir(cp_gf,  $\beta$ ,  $\gamma$ )
```

```
EdgeModelSystem(Model clustered_sir:
```

```
Equations (8):
```

```
8 standard: see equations(clustered_sir)
```

```
Unknowns (8): see unknowns(clustered_sir)
```

```
pop_R(t)
```

```
pop_I(t)
```

```
 $\varphi_3_R(t)$ 
```

```
 $\varphi_3_I(t)$ 
```

```
 $\varphi_2$ 
```

```
Parameters (3): see parameters(clustered_sir)
```

```
 $\rho$ 
```

```
 $\beta$ 
```

```
 $\gamma$ 
```

```
Observed (4): see observed(clustered_sir), Dict{Symbol, Any}(: $\theta_3 \Rightarrow \theta_3(t)$ , : $\varphi_3_I \Rightarrow \varphi_3_I(t)$ , :  
1 + 3.0(-1 +  $\theta_2(t)$ ) +  $\theta_3(t)^2$ )), :rho_param =>  $\rho$ , :edge_seed_groups => Any[(entry =  $\varphi_2_I(t)$ , t  
1 + 3.0(-1 +  $\theta_2(t)$ ) +  $\theta_3(t)^2$ ) - exp(-1 + 3.0(-1 +  $\theta_2(t)$ ) +  $\theta_3(t)^2$ )* $\rho$ ) / exp(0.0)), (entry =  $\varphi_3_I$   
1 + 3.0(-1 +  $\theta_2(t)$ ) +  $\theta_3(t)^2$ )*(1 -  $\rho$ ) / exp(0.0))]]))
```

The clustered model tracks more state variables due to the separate single-edge and triangle-edge dynamics:

```
println("Unclustered variables: ", keys(model_unclustered.variables))
println("Clustered variables:   ", keys(model_clustered.variables))
```

```
Unclustered variables: [:R, : $\varphi_I$ , :pop_I, :pop_R, : $\varphi_R$ , : $\theta$ ]
```

```
Clustered variables:  [: $\theta_3$ , : $\varphi_3_I$ , : $\varphi_3_R$ , : $\varphi_2_I$ , :pop_I, :R, :pop_R, : $\varphi_2_R$ , : $\theta_2$ ]
```

Comparing clustered vs unclustered epidemics

We now solve both models with the unclustered Poisson baseline anchored at $R_0 = 2$ ($\beta = 1/6$, $\gamma = 0.25$), and compare the epidemic curves.

```
tspan = (0.0, 40.0)
params = Dict( $\beta \Rightarrow \beta_{val}$ ,  $\gamma \Rightarrow \gamma_{val}$ )
```

```

# Unclustered
ic1 = merge(default_initial_conditions(model_unclustered), params)
prob1 = ODEProblem(model_unclustered.system, ic1, tspan)
sol1 = solve(prob1, Tsit5(); saveat = 0.5)

# Clustered
ic2 = merge(default_initial_conditions(model_clustered), params)
prob2 = ODEProblem(model_clustered.system, ic2, tspan)
sol2 = solve(prob2, Tsit5(); saveat = 0.5)

```

retcode: Success

Interpolation: 1st order linear

t: 81-element Vector{Float64}:

```

0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0
4.5
⋮
36.0
36.5
37.0
37.5
38.0
38.5
39.0
39.5
40.0

```

u: 81-element Vector{Vector{Float64}}:

```

[0.0, 0.001, 0.0, 0.0010000000000000009, 0.0, 0.0010000000000000009, 1.0, 1.0]
[0.000144623672614379, 0.001326107805807947, 0.00014461643846619562, 0.001325973247055298,
[0.00033497795415506396, 0.0017352069162356533, 0.0003349379878670594, 0.00173478462655592
[0.0005828644053464031, 0.0022510324558243327, 0.0005827405493116342, 0.002250055298484182
[0.0009034147604395979, 0.0029034455012692355, 0.0009031116558782131, 0.002901456725756914
[0.0013159697685541298, 0.0037299981287621455, 0.0013153165442927257, 0.003726225967734569
[0.001845104935954855, 0.004777700564948291, 0.0018438057262388148, 0.004770859951768346, 0
[0.0025220405153348123, 0.00610529737504897, 0.0025195880956793057, 0.0060932557069970986,

```

```
[0.0033860764081957277, 0.007785467336545977, 0.0033816217012115767, 0.007764739705640867,
[0.004486804708450796, 0.009907869260592578, 0.004478935779073265, 0.009872758398233766, 0.
[0.7684924194637379, 0.007192733932296931, 0.5103843458746748, 0.0013664247117114415, 0.464
[0.7693445792642608, 0.006452802328367567, 0.5105445393464301, 0.0011974197922827636, 0.464
[0.7701087548816637, 0.005786927653077036, 0.5106847513725135, 0.0010495222828236022, 0.464
[0.7707938267003962, 0.005187961768522474, 0.5108074749911291, 0.0009201003717910553, 0.464
[0.7714078772463105, 0.004649395858668374, 0.5109149871101936, 0.0008067432479673713, 0.465
[0.7719581911297007, 0.004165360386799442, 0.5110093483369407, 0.0007072612939204892, 0.465
[0.7724512024872386, 0.003730549318561273, 0.5110921803987606, 0.0006199402285695773, 0.465
[0.7728926785078177, 0.0033401151108894125, 0.511164784159956, 0.0005434150928230247, 0.465
[0.7732878931740367, 0.0029896659767265237, 0.5112284260589087, 0.00047634766958477864, 0.4
```

Extract the epidemic compartments using the PGFs:

```
S1 = compartment(sol1, model_unclustered, :S)
I1 = compartment(sol1, model_unclustered, :I)
R1 = compartment(sol1, model_unclustered, :R)

S2 = compartment(sol2, model_clustered, :S)
I2 = compartment(sol2, model_clustered, :I)
R2 = compartment(sol2, model_clustered, :R)
```

81-element Vector{Float64}:

```
0.0
0.000144623672614379
0.00033497795415506396
0.0005828644053464031
0.0009034147604395979
0.0013159697685541298
0.001845104935954855
0.0025220405153348123
0.0033860764081957277
0.004486804708450796
0.7684924194637379
0.7693445792642608
0.7701087548816637
0.7707938267003962
0.7714078772463105
0.7719581911297007
0.7724512024872386
```

0.7728926785078177
0.7732878931740367

```
plot(sol1.t, S1, label = "S (unclustered)", linewidth = 2, color = :blue)
plot!(sol1.t, I1, label = "I (unclustered)", linewidth = 2, color = :red)
plot!(sol1.t, R1, label = "R (unclustered)", linewidth = 2, color = :green)
plot!(sol2.t, S2, label = "S (clustered)", linewidth = 2, color = :blue, linestyle = :dash)
plot!(sol2.t, I2, label = "I (clustered)", linewidth = 2, color = :red, linestyle = :dash)
plot!(sol2.t, R2, label = "R (clustered)", linewidth = 2, color = :green, linestyle = :dash)
xlabel!("Time")
ylabel!("Fraction of population")
title!("Clustered vs Unclustered SIR ( $R_0=2$  unclustered baseline)")
```




Figure 1: Figure 1: Clustering reduces peak prevalence and delays the epidemic peak. Both networks have mean degree 5.

The dashed (clustered) epidemic has a lower peak prevalence and a delayed peak compared to the solid (unclustered) curves. The final epidemic size is also smaller with clustering.

Varying the clustering coefficient

We fix the total mean degree at 5 and vary the fraction of edges that come from triangles. As clustering increases, more edges are “wasted” on redundant triangle paths.

```
fracs = [0.0, 0.2, 0.4, 0.6]
results = []

for frac in fracs
    κ_s_i = 5.0 * (1 - frac)
    κ_t_i = 5.0 * frac / 2 # each triangle stub contributes 2 neighbours
    cpgf_i = clustered_poisson_pgf(κ_s_i, κ_t_i)
    model_i = build_clustered_sir(cpgf_i, β, γ)
    ic_i = merge(default_initial_conditions(model_i), params)
    prob_i = ODEProblem(model_i.system, ic_i, tspan)
    sol_i = solve(prob_i, Tsit5(); saveat = 0.5)

    S_i = compartment(sol_i, model_i, :S)
    I_i = compartment(sol_i, model_i, :I)
    R_i_vals = compartment(sol_i, model_i, :R)
    C_i = 2κ_t_i / (2κ_t_i + κ_s_i)

    push!(results, (frac = frac, C = C_i, t = sol_i.t, I = I_i, S = S_i, R = R_i_vals))
end

p = plot(xlabel = "Time", ylabel = "Fraction infected",
         title = "Effect of clustering on epidemic dynamics")
colors = [:red, :orange, :purple, :blue]
for (i, res) in enumerate(results)
    plot!(p, res.t, res.I,
          label = "C = $(round(res.C; digits=2)) (frac = $(res.frac))",
          linewidth = 2, color = colors[i])
end
p
```

As the clustering coefficient C increases, the peak prevalence decreases and the epidemic is delayed.

Clustering coefficient and R_0

The basic reproduction number R_0 for a clustered network accounts for the redundant transmission paths within triangles. We can compute it using `basic_reproduction_number`:




Figure 2: Figure 2: Increasing the fraction of triangle edges reduces epidemic severity while keeping total mean degree = 5.

```

using EdgeBasedModels: ClusteredConfigurationModel, DiseaseProgression, DiseaseStage, Dise

fracs_fine = 0.0:0.05:0.8
C_vals = Float64[]
R0_vals = Float64[]

for frac in fracs_fine
    κ_s_i = 5.0 * (1 - frac)
    κ_t_i = 5.0 * frac / 2
    cp_gf_i = clustered_poisson_pgf(κ_s_i, κ_t_i)
    C_i = 2κ_t_i / (2κ_t_i + κ_s_i)
    push!(C_vals, C_i)

    prog = DiseaseProgression(
        [DiseaseStage(:I; transmission_rate = β_val),
         DiseaseStage(:R; transmission_rate = 0)],
        [DiseaseTransition(:I, :R, γ_val)]; entry = :I)
    cm = ClusteredConfigurationModel(cp_gf_i, prog)
    R0_i = basic_reproduction_number(cm)
    push!(R0_vals, Float64(Symbolics.value(R0_i)))
end

plot(C_vals, R0_vals, linewidth = 2, color = :darkred, marker = :circle, markersize = 3,
      label = "R0")
xlabel!("Clustering coefficient C")
ylabel!("R0")
title!("R0 vs Clustering (total mean degree = 5, anchored β, γ)")
hline!([1.0], linestyle = :dash, color = :gray, label = "R0 = 1")

```

As the clustering coefficient rises from 0 toward 1, R_0 decreases. This quantifies the protective effect of clustering: triangles create correlated transmission paths that reduce the effective number of secondary infections.

SEIR with clustering

The clustered EBCM framework extends to arbitrary disease progressions. Here we add a latent (exposed) period using `build_clustered_seir`:

```

@parameters σ

cp_gf_seir = clustered_poisson_pgf(3.0, 1.0)

```




Figure 3: R_0 decreases as the clustering coefficient increases, reflecting the redundancy of triangle edges.

```
model_seir = build_clustered_seir(cpgf_seir,  $\sigma$ ,  $\beta$ ,  $\gamma$ )
println("SEIR clustered variables: ", keys(model_seir.variables))
```

SEIR clustered variables: [: θ_3 , : φ_3_I , : φ_3_R , : φ_2_I , :pop_I, :R, : θ_2 , : φ_2_E , : φ_3_E , :pop_R, : φ_1]

```
params_seir = Dict( $\beta \Rightarrow \beta_{\text{val}}$ ,  $\gamma \Rightarrow \gamma_{\text{val}}$ ,  $\sigma \Rightarrow 0.2$ )
ic_seir = merge(default_initial_conditions(model_seir), params_seir)
prob_seir = ODEProblem(model_seir.system, ic_seir, tspan)
sol_seir = solve(prob_seir, Tsit5(); saveat = 0.5)

S_seir = compartment(sol_seir, model_seir, :S)
R_seir_vals = compartment(sol_seir, model_seir, :R)
I_seir = 1.0 .- S_seir .- R_seir_vals
```

81-element Vector{Float64}:

```
0.00100000000000000009
0.0010130454905676778
0.00104717512584005
0.0010968014650383033
0.0011582546048606796
0.0012291684860312529
0.0013080987956067896
0.0013942005315523895
0.0014870910042569122
0.0015866444132882913
␣
0.07593289690599012
0.07978103926137754
0.08374450570741221
0.08781939417069416
0.09200168706017353
0.09628717125938369
0.10066838862572111
0.10513033548489154
0.10966182705534845
```

```
plot(sol_seir.t, S_seir, label = "S", linewidth = 2, color = :blue)
plot!(sol_seir.t, I_seir, label = "E + I", linewidth = 2, color = :red)
plot!(sol_seir.t, R_seir_vals, label = "R", linewidth = 2, color = :green)
xlabel!("Time")
```

```
ylabel!("Fraction of population")
title!("Clustered SEIR ( $\kappa_s=3$ ,  $\kappa_t=1$ ,  $\sigma=0.2$ , anchored  $\beta$ ,  $\gamma$ )")
```




Figure 4: SEIR epidemic on a clustered network with a latent period ($\sigma=0.2$).

The latent period further delays and flattens the epidemic curve on top of the clustering effect.

Summary

- Clustering (triadic closure) is a key structural property of real contact networks that is absent from the standard configuration model.
- Triangles create redundant transmission paths: if two of three triangle edges transmit, the third cannot cause a new infection.
- The clustered EBCM uses a bivariate PGF $g(x, y)$ and tracks separate survival probabilities for single-edge (θ_2) and triangle-edge (θ_3) partners.
- Increasing clustering reduces R_0 and lowers epidemic peak prevalence while keeping the mean degree fixed.

- The `EdgeBasedModels.jl` package provides `clustered_poisson_pgf`, `build_clustered_sir`, and `build_clustered_seir` for modelling clustered networks with arbitrary disease progressions.

Simulation validation

This vignette focuses on a scenario for which `NetworkOutbreaks.jl` does not yet provide an out-of-the-box stochastic ground truth (multi-type host graphs with prescribed mixing matrices, time-varying networks via the EBCM rewiring schedule, multiplex layers, degree-correlated configuration models, clustered graphs with prescribed triangle counts, or final-size sweeps over R_0).

A future revision will add the missing primitives to `NetworkOutbreaks.jl` and overlay the corresponding Gillespie SSA ribbon here. Until then, see [vignette 01](#) for the validation pattern on a single-layer Poisson configuration-model network with the same canonical parameters.