

Multiplex Networks

EdgeBasedModels.jl Vignette 11

Simon Frost

2026-05-14

Table of contents

Introduction	1
Setup	2
Two-Layer Model	3
Layer dynamics	5
Population-level epidemic	6
Layer Contributions	7
Three-Layer Example	10
Varying Layer Structure	13
R_0 Decomposition	16
Intervention targeting	17
Summary	19
Simulation validation	20

Introduction

Real populations do not interact through a single homogeneous contact network. Instead, individuals maintain several distinct contact layers — households, workplaces, schools, community settings — each with its own degree distribution and transmission characteristics. A household layer may be dense (few contacts, high transmission probability), while a community layer may be sparse (many weak contacts).

The multiplex EBCM extends the edge-based compartmental model to handle this structure. Each layer i has its own probability generating function ψ_i , transmission rate β_i , and recovery rate γ_i . The key insight is that layers are independent conditional on a node's infection status, so each layer gets its own edge-level variable $\theta_i(t)$. The population-level susceptible fraction is then:

$$S(t) = \prod_i \psi_i(\theta_i(t))$$

Each θ_i evolves according to the standard EBCM equations for its layer, but the coupling arises because the infection of a node through *any* layer removes it from the susceptible pool across *all* layers.

This approach was described by Jacobsen et al. (2016) and builds on the single-layer framework of Volz (2008) and Miller, Slim & Volz (2012).

Setup

```
using EdgeBasedModels
using ModelingToolkit
using OrdinaryDiffEq
using Plots
```

We define helper functions to extract trajectories from multiplex ODE solutions.

```
"""
Extract a named trajectory from a multiplex solution.
"""
function named_trajectory(sol, sys, name::Symbol)
    target = string(name) * "(t)"
    var = findfirst(v → string(v) == target || endswith(string(v), target), ModelingToolkit.
    var == nothing && throw(ArgumentError("unknown trajectory: $name"))
    return sol[ModelingToolkit.unknowns(sys)[var]]
end

"""
Extract θ values from a multiplex solution by matching variable names.
"""
function extract_θ(sol, sys, layer_names)
    θ_vals = Dict{Symbol, Vector{Float64}}{ }
    for name in layer_names
        θ_vals[name] = named_trajectory(sol, sys, Symbol("θ_$name"))
    end
    return θ_vals
end
```

```

"""
Extract the R trajectory from a multiplex solution.
"""
function extract_R(sol, sys)
    return named_trajectory(sol, sys, :R)
end

"""
Compute  $S(t) = \mathbb{I}\Psi(\theta)$  for Poisson layers with given mean degrees.
"""
function compute_S( $\theta$ _dict, mean_degrees)
    layer_names = collect(keys( $\theta$ _dict))
    n = length(first(values( $\theta$ _dict)))
    S = ones(n)
    for name in layer_names
         $\kappa$  = mean_degrees[name]
        S .*= exp( $\kappa$  .* ( $\theta$ _dict[name] .- 1))
    end
    return S
end

```

Main.Notebook.compute_S

Two-Layer Model

We start with two contact layers representing distinct social settings:

- Household: dense contacts (Poisson mean degree 3), per-edge transmissibility $T = 0.75$ ($\beta = 0.75$), recovery rate $\gamma = 0.25$.
- Community: more contacts (Poisson mean degree 8), per-edge transmissibility $T = 0.5$ ($\beta = 0.25$), same recovery rate.

```

pgf_hh = poisson_pgf(3.0)
pgf_cm = poisson_pgf(8.0)

layers_2 = [
    (:household, pgf_hh, 0.75, 0.25),
    (:community, pgf_cm, 0.25, 0.25),
]

(sys_2, u0_2, tspan_2, p_2) = build_multiplex_sir(layers_2; tspan = (0.0, 40.0))

```

```

(Model multiplex_sir:
Equations (3):
  3 standard: see equations(multiplex_sir)
Unknowns (3): see unknowns(multiplex_sir)
  R(t)
   $\theta_{\text{community}}(t)$ 
   $\theta_{\text{household}}(t)$ 
Parameters (4): see parameters(multiplex_sir)
   $\beta_{\text{household}}$ 
   $\gamma_{\text{household}}$ 
   $\gamma_{\text{community}}$ 
   $\beta_{\text{community}}$ , Dict{Any, Float64}( $\theta_{\text{community}}(t) \Rightarrow 0.999999$ ,  $\theta_{\text{household}}(t) \Rightarrow 0.999999$ , R(t)

```

```

prob_2 = ODEProblem(sys_2, merge(u0_2, p_2), tspan_2)
sol_2 = solve(prob_2, Tsit5(); saveat = 0.5)

```

```

retcode: Success
Interpolation: 1st order linear
t: 81-element Vector{Float64}:
  0.0
  0.5
  1.0
  1.5
  2.0
  2.5
  3.0
  3.5
  4.0
  4.5
  ⋮
 36.0
 36.5
 37.0
 37.5
 38.0
 38.5
 39.0
 39.5
 40.0
u: 81-element Vector{Vector{Float64}}:
 [0.0, 0.999999, 0.999999]
 [3.63279446628779e-6, 0.999995742264062, 0.9999897915403911]

```

```
[2.422259667017793e-5, 0.9999765605470764, 0.9999382828010519]
[0.00014213475595432605, 0.9998661230469152, 0.999643560035654]
[0.0008158107420174954, 0.999234796551605, 0.9979606301411036]
[0.004546433575994476, 0.9957424583906558, 0.9886719572364643]
[0.022311257119507554, 0.979211827512088, 0.9451808429462397]
[0.07572686335067384, 0.9306173065698817, 0.8224450832543508]
[0.16023638595115797, 0.8577967265479756, 0.6557695976706662]
[0.2506693936188328, 0.7863253926489508, 0.5161292964972096]
```

□

```
[0.997756473556317, 0.5009767458991977, 0.2516374960942665]
[0.9977905840619087, 0.500976313043348, 0.2517821580952215]
[0.997820466090964, 0.5009764647141596, 0.251721918244749]
[0.9978466768797258, 0.5009769516341833, 0.25154493193008315]
[0.9978698572079168, 0.5009773098758451, 0.2514139455614775]
[0.997890472508553, 0.5009773717545006, 0.25138705707824005]
[0.99790884218409, 0.5009771438441062, 0.25146188562276606]
[0.997925139606422, 0.5009768069772188, 0.2515755715405377]
[0.997939392116882, 0.5009767162449962, 0.25160477638012474]
```

Layer dynamics

Each layer has its own $\theta_i(t)$, representing the probability that a random edge in layer i has not yet transmitted infection. These evolve at different rates depending on the layer's transmission parameters.

```
layer_names_2 = [:household, :community]
θ_2 = extract_θ(sol_2, sys_2, layer_names_2)

plot(sol_2.t, θ_2[:household], label = "θ (household)", linewidth = 2, color = :red)
plot!(sol_2.t, θ_2[:community], label = "θ (community)", linewidth = 2, color = :blue)
xlabel!("Time")
ylabel!("θ(t)")
title!("Edge-level dynamics by layer")
```




Figure 1: Edge-level probabilities θ for household and community layers. The household layer (higher β) drops faster despite having fewer contacts.

Population-level epidemic

The susceptible fraction is the product $S(t) = \psi_{hh}(\theta_{hh}) \cdot \psi_{cm}(\theta_{cm})$.

```
mean_degrees_2 = Dict(:household => 3.0, :community => 8.0)
S_2 = compute_S(theta_2, mean_degrees_2)
R_2 = extract_R(sol_2, sys_2)
I_2 = 1.0 .- S_2 .- R_2

plot(sol_2.t, S_2, label = "S(t)", linewidth = 2, color = :green)
plot!(sol_2.t, I_2, label = "I(t)", linewidth = 2, color = :red)
plot!(sol_2.t, R_2, label = "R(t)", linewidth = 2, color = :blue)
xlabel!("Time")
ylabel!("Fraction of population")
title!("Two-layer multiplex epidemic")
```




Figure 2: Two-layer multiplex SIR epidemic. S, I, and R trajectories computed from per-layer θ values and the shared recovery variable.

Layer Contributions

Which layer drives the epidemic? We can answer this by comparing the full two-layer model to single-layer models where one layer is “switched off” (by setting its $\beta = 0$).

```
layers_hh_only = [
    (:household, pgf_hh, 0.75, 0.25),
    (:community, pgf_cm, 0.0, 0.25), # community switched off
]
layers_cm_only = [
    (:household, pgf_hh, 0.0, 0.25), # household switched off
    (:community, pgf_cm, 0.25, 0.25),
]

(sys_hh, u0_hh, ts_hh, p_hh) = build_multiplex_sir(layers_hh_only; tspan = (0.0, 40.0))
(sys_cm, u0_cm, ts_cm, p_cm) = build_multiplex_sir(layers_cm_only; tspan = (0.0, 40.0))

sol_hh = solve(ODEProblem(sys_hh, merge(u0_hh, p_hh), ts_hh), Tsit5(); saveat = 0.5)
sol_cm = solve(ODEProblem(sys_cm, merge(u0_cm, p_cm), ts_cm), Tsit5(); saveat = 0.5)
```

retcode: Success

Interpolation: 1st order linear

t: 81-element Vector{Float64}:

0.0

0.5

1.0

1.5

2.0

2.5

3.0

3.5

4.0

4.5

⋮

36.0

36.5

37.0

37.5

38.0

38.5

39.0

39.5

40.0

u: 81-element Vector{Vector{Float64}}:

[0.0, 0.999999, 0.999999]

[1.9682069851703285e-6, 0.9999973539570957, 0.9999991175031212]

[5.960822364933464e-6, 0.9999939324789047, 0.9999992211992138]

[1.4244389589973176e-5, 0.9999867522749707, 0.9999993127107802]

[3.1654318794299666e-5, 0.9999715824689565, 0.999999393469343]

[6.827212703831255e-5, 0.9999396022922835, 0.9999994647386083]

[0.00014581210139864354, 0.9998718143451505, 0.9999995276334621]

[0.00030920178586646066, 0.9997289171760659, 0.9999995831379876]

[0.0006549410705557083, 0.9994265237111511, 0.9999996321205747]

[0.0013819926810231594, 0.9987907420501887, 0.9999996753475187]

⋮

[0.9792306230573605, 0.5099140076194297, 0.999999998775039]

[0.9793402265924366, 0.5099140910202603, 0.99999999891754]

[0.9794355347944357, 0.5099145179106217, 0.999999999041369]

[0.9795186881804505, 0.5099151325160176, 0.999999999149347]

[0.9795918648460826, 0.5099157223677209, 0.999999999244357]

[0.9796572804654428, 0.5099160183027741, 0.999999999329348]

[0.979717052162368, 0.5099157333625755, 0.99999999940714]

[0.9797705466170944, 0.5099153154117279, 0.999999999476803]


```
[0.97981776202141, 0.5099149834325335, 0.999999999953828]
```

```
theta_hh_only = extract_theta(sol_hh, sys_hh, layer_names_2)
theta_cm_only = extract_theta(sol_cm, sys_cm, layer_names_2)

S_hh_only = compute_S(theta_hh_only, mean_degrees_2)
S_cm_only = compute_S(theta_cm_only, mean_degrees_2)
R_hh_only = extract_R(sol_hh, sys_hh)
R_cm_only = extract_R(sol_cm, sys_cm)

plot(sol_2.t, R_2, label = "Both layers", linewidth = 2.5, color = :black)
plot!(sol_hh.t, R_hh_only, label = "Household only", linewidth = 2,
      color = :red, linestyle = :dash)
plot!(sol_cm.t, R_cm_only, label = "Community only", linewidth = 2,
      color = :blue, linestyle = :dash)
xlabel!("Time")
ylabel!("Cumulative recovered fraction")
title!("Layer contributions to the epidemic")
```




Figure 3: Epidemic curves with individual layers vs both layers active. The combined epidemic is larger than either layer alone, since R_0 is additive across layers.

The total R_0 for a multiplex model is the sum of per-layer contributions:

```

R0_total = multiplex_R0(layers_2)
R0_hh = multiplex_R0(layers_hh_only)
R0_cm = multiplex_R0(layers_cm_only)
println("R0 (both layers):      $(round(R0_total; digits = 3))")
println("R0 (household only):   $(round(R0_hh; digits = 3))")
println("R0 (community only):    $(round(R0_cm; digits = 3))")
println("Sum of single layers:  $(round(R0_hh + R0_cm; digits = 3))")

```

```

R0 (both layers):      6.25
R0 (household only):   2.25
R0 (community only):   4.0
Sum of single layers:  6.25

```

Three-Layer Example

Adding more layers is straightforward. In the compact multiplex implementation used here, each additional layer introduces one additional edge-survival equation θ_i , together with the shared recovery equation R . So an L -layer model has $L + 1$ ODEs in total.

Consider a population with household, school, and community contacts:

- Household: Poisson(3), $T = 0.8$ ($\beta = 1.0$), $\gamma = 0.25$ — few but intense contacts.
- School: Poisson(15), $T = 1/3$ ($\beta = 0.125$), $\gamma = 0.25$ — many contacts, low per-contact transmission.
- Community: Poisson(5), $T = 0.6$ ($\beta = 0.375$), $\gamma = 0.25$ — moderate in both dimensions.

```

pgf_school = poisson_pgf(15.0)
pgf_cm3 = poisson_pgf(5.0)
pgf_hh3 = poisson_pgf(3.0)

layers_3 = [
    (:household, pgf_hh3, 1.0, 0.25),
    (:school, pgf_school, 0.125, 0.25),
    (:community, pgf_cm3, 0.375, 0.25),
]

(sys_3, u0_3, ts_3, p_3) = build_multiplex_sir(layers_3; tspan = (0.0, 40.0))
sol_3 = solve(ODEProblem(sys_3, merge(u0_3, p_3), ts_3), Tsit5(); saveat = 0.5)

```

```

retcode: Success
Interpolation: 1st order linear
t: 81-element Vector{Float64}:
 0.0
 0.5
 1.0
 1.5
 2.0
 2.5
 3.0
 3.5
 4.0
 4.5
 ⋮
36.0
36.5
37.0
37.5
38.0
38.5
39.0
39.5
40.0
u: 81-element Vector{Vector{Float64}}:
 [0.0, 0.999999, 0.999999, 0.999999]
 [1.783810622744761e-5, 0.9999738293319783, 0.9999904026667394, 0.9999368546517261]
 [0.00036165574780582514, 0.9994876881811198, 0.9998220269174593, 0.9987529907087456]
 [0.006487649656226426, 0.9908470081169506, 0.9968225207246775, 0.9778010042627571]
 [0.057432350354192675, 0.9208477192841362, 0.9720990503295311, 0.81473544654546]
 [0.15585921429416905, 0.7974321028858508, 0.9257960497877081, 0.5648926152428037]
 [0.25329408669449993, 0.6931119100724087, 0.8823363202573526, 0.40048708405130645]
 [0.34063657104773704, 0.6149797281028747, 0.8456547921111012, 0.3084937698010898]
 [0.4179777923194391, 0.5574665291676003, 0.8151187083133106, 0.2584953174724962]
 [0.4863054196310205, 0.515288933556277, 0.7897676005269574, 0.2315047862852041]
 ⋮
 [0.9997740028941887, 0.40001827298798553, 0.6666777551677318, 0.19992754090982262]
 [0.999796979079336, 0.40001826884456937, 0.6666775938090409, 0.19999323313184483]
 [0.9998172545940696, 0.4000182609174536, 0.6666774591857255, 0.20013218249756473]
 [0.9998351490716251, 0.40001826013125386, 0.666677348722546, 0.20014089011600394]
 [0.9998509439627565, 0.40001826596435003, 0.666677258367058, 0.20002838298341405]
 [0.9998648825118841, 0.4000182692878587, 0.6666771831213587, 0.19996321735418415]
 [0.999877181371112, 0.40001826764499177, 0.6666771198453879, 0.1999905053746088]
 [0.9998880343962171, 0.4000182628616879, 0.6666770666848597, 0.20007644320698464]

```

```
[0.9998976126466488, 0.4000182610466125, 0.6666770230712625, 0.20010831102961008]
```

```
layer_names_3 = [:household, :school, :community]
θ_3 = extract_θ(sol_3, sys_3, layer_names_3)

plot(sol_3.t, θ_3[:household], label = "θ (household)", linewidth = 2, color = :red)
plot!(sol_3.t, θ_3[:school], label = "θ (school)", linewidth = 2, color = :orange)
plot!(sol_3.t, θ_3[:community], label = "θ (community)", linewidth = 2, color = :blue)
xlabel!("Time")
ylabel!("θ(t)")
title!("Three-layer edge dynamics")
```




Figure 4: Per-layer θ dynamics in the three-layer model. Each layer decays at its own rate, reflecting different transmission characteristics.

```
mean_degrees_3 = Dict{:household ⇒ 3.0, :school ⇒ 15.0, :community ⇒ 5.0}
S_3 = compute_S(θ_3, mean_degrees_3)
R_3 = extract_R(sol_3, sys_3)
I_3 = 1.0 .- S_3 .- R_3

plot(sol_3.t, S_3, label = "S(t)", linewidth = 2, color = :green)
plot!(sol_3.t, I_3, label = "I(t)", linewidth = 2, color = :red)
```

```

plot(sol_3.t, R_3, label = "R(t)", linewidth = 2, color = :blue)
xlabel!("Time")
ylabel!("Fraction of population")
title!("Three-layer multiplex epidemic")

```




Figure 5: Three-layer multiplex epidemic. The product formula $S(t) = \prod \psi_i(\theta_i)$ combines contributions from all layers.

Despite having the highest mean degree, the school layer has the lowest per-contact transmission rate. The household layer, with few contacts but high β , may contribute more to overall transmission. We quantify this next.

Varying Layer Structure

How does the distribution of contacts across layers affect epidemic outcomes? We fix the total mean contacts at 10 and vary the allocation between a “close” layer ($T = 0.75$, $\beta = 0.75$) and a “casual” layer ($T = 4/9 \approx 0.44$, $\beta = 0.2$), both with $\gamma = 0.25$.

```

close_frcs = 0.05:0.05:0.95
final_infected = Float64[]
final_theta_close = Float64[]

```

```

final_theta_casual = Float64[]

total_contacts = 10.0
beta_close = 0.75
beta_casual = 0.2
gamma_val = 0.25

for f in close_fracs
    k_close = f * total_contacts
    k_casual = (1 - f) * total_contacts
    pgf_c = poisson_pgf(k_close)
    pgf_a = poisson_pgf(k_casual)

    layers_v = [
        (:close, pgf_c, beta_close, gamma_val),
        (:casual, pgf_a, beta_casual, gamma_val),
    ]

    (sys_v, u0_v, ts_v, p_v) = build_multiplex_sir(layers_v; tspan = (0.0, 80.0))
    sol_v = solve(ODEProblem(sys_v, merge(u0_v, p_v), ts_v), Tsit5())

    R_v = extract_R(sol_v, sys_v)
    theta_v = extract_theta(sol_v, sys_v, [:close, :casual])
    theta_c_end = theta_v[:close][end]
    theta_a_end = theta_v[:casual][end]
    push!(final_infected, R_v[end])
    push!(final_theta_close, theta_c_end)
    push!(final_theta_casual, theta_a_end)
end

plot(collect(close_fracs), final_infected,
     label = "Final epidemic size", linewidth = 2, color = :black,
     marker = :circle, markersize = 3)
xlabel!("Fraction of contacts in close layer")
ylabel!("Final infected fraction")
title!("Contact allocation (total mean = $total_contacts)")

```




Figure 6: Effect of redistributing fixed total contacts between a high- β (close) and low- β (casual) layer. The final epidemic size ($1 - S(\infty)$) varies with the allocation.

```
plot(collect(close_fracs), final_theta_close,
      label = "θ∞ (close layer)", linewidth = 2, color = :red)
plot!(collect(close_fracs), final_theta_casual,
        label = "θ∞ (casual layer)", linewidth = 2, color = :blue)
xlabel!("Fraction of contacts in close layer")
ylabel!("Final θ value")
title!("Per-layer edge transmission")
```




Figure 7: Final θ values for each layer as the contact allocation changes. The close layer (high β) shows more transmission per edge when it has fewer contacts.

The final epidemic size depends on how contacts are partitioned. Even with the same total contacts, concentrating more contacts in the high- β layer can increase the overall epidemic size — the nonlinear coupling between layers means that the distribution matters, not just the total.

R Decomposition

For a multiplex model, $R_0 = \sum_i T_i \cdot \langle k_i^2 - k_i \rangle / \langle k_i \rangle$, where $T_i = \beta_i / (\beta_i + \gamma_i)$ is the transmissibility of layer i . This additive structure makes it easy to identify which layer contributes most — and where interventions would be most effective.

```
function layer_R0(name, pgf,  $\beta$ ,  $\gamma$ )
  T =  $\beta$  / ( $\beta$  +  $\gamma$ )
   $\psi\_fn$  = x → EdgeBasedModels._build_pgf_fn(pgf)(x)
   $\psi1\_fn$  = x → EdgeBasedModels._build_pgf_deriv_fn(pgf, 1)(x)
   $\psi2\_fn$  = x → EdgeBasedModels._build_pgf_deriv_fn(pgf, 2)(x)
  mean_k =  $\psi1\_fn(1.0)$ 
  excess =  $\psi2\_fn(1.0)$  / mean_k
  return T * excess
end
```



```

r0_contributions = [(String(l[1]), layer_R0(l...)) for l in layers_3]
names_bar = [c[1] for c in r0_contributions]
vals_bar = [c[2] for c in r0_contributions]

bar(names_bar, vals_bar,
    label = "Layer  $R_0$ ",
    color = [:red, :orange, :blue],
    ylabel = " $R_0$  contribution",
    title = " $R_0$  decomposition (total = $(round(sum(vals_bar); digits = 2)))",
    legend = false)

```




Figure 8: Per-layer R_0 contributions for the three-layer model. Targeting the layer with the largest R_0 contribution yields the greatest reduction.

Intervention targeting

Suppose we can reduce β in one layer by 50%. Which layer should we target?

```

interventions = Dict{String,Float64}{}

for (i, (name, pgf_l,  $\beta_l$ ,  $\gamma_l$ )) in enumerate(layers_3)

```

```

modified = collect(layers_3)
modified[i] = (name, pgf_l,  $\beta_l$  * 0.5,  $\gamma_l$ )

(sys_int, u0_int, ts_int, p_int) = build_multiplex_sir(modified; tspan = (0.0, 40.0))
sol_int = solve(ODEProblem(sys_int, merge(u0_int, p_int), ts_int), Tsit5())

R_int = extract_R(sol_int, sys_int)
interventions[String(name)] = R_int[end]
end

# Baseline
baseline_final = R_3[end]

bar_names = ["Baseline"; collect(keys(interventions))]
bar_vals = [baseline_final; collect(values(interventions))]
bar_colors = [:gray, :red, :orange, :blue]

bar(bar_names, bar_vals,
    label = false,
    color = bar_colors,
    ylabel = "Final recovered fraction",
    title = "50%  $\beta$  reduction – which layer to target?")

```




Figure 9: Effect of a 50% reduction in β on each layer. Targeting the layer with the highest R_0 contribution produces the largest reduction in epidemic size.

The layer with the largest R_0 contribution is the most impactful intervention target. This decomposition provides a principled basis for allocating limited resources across contact settings.

Summary

- Multiplex EBCM extends the edge-based framework to populations with multiple independent contact layers, each with its own degree distribution and transmission dynamics.
- Each layer has its own $\theta_i(t)$; the susceptible fraction is $S(t) = \prod_i \psi_i(\theta_i)$.
- `build_multiplex_sir(layers)` constructs the ODE system from a vector of (name, pgf, β , γ) tuples.
- `multiplex_R0(layers)` computes the total $R_0 = \sum_i R_{0,i}$, with each layer contributing independently.
- `susceptible_fraction(pgfs,  $\theta$ _values)` evaluates the product formula at any time point.
- Contact allocation matters: redistributing fixed total contacts across layers with different transmission rates changes epidemic outcomes.
- Intervention targeting: the per-layer R_0 decomposition identifies which layer to prioritise for control measures.

Simulation validation

We validate the two-layer ODE prediction against a stochastic Gillespie SSA on `MultiplexNetwork` from `NetworkOutbreaks.jl`. Each layer is sampled as an independent $\text{Poisson}(\kappa_i)$ Erdős-Rényi graph; the per-layer transmission rate β_i enters as the layer's `layer_rate` weight on a unit-rate infection transition.

Note: the default initial conditions returned by `build_multiplex_sir` correspond to an infinitesimal seed ($\theta_i(0) = 1 - 10^{-6}$). For a fair comparison with a finite-population SSA seeded at fraction f , we re-solve the multiplex ODE with matched ICs: each layer's $\theta_i(0) = \psi_i^{-1}((1 - f)^{1/n})$ so that $S(0) = \prod \psi_i(\theta_i) = 1 - f$. For $\text{Poisson}(\kappa)$ PGFs the inverse is $\theta = 1 + \log((1-f)^{1/n}) / \kappa$.

```
include("../_validation.jl")

const N_sim = 1000
const seed_fraction_val = 0.01
const tgrid_val = collect(0.0:0.5:40.0)

# --- Re-solve the EBCM with seed-fraction-matched ICs -----
# For Poisson( $\kappa$ ) PGF  $\psi(x) = \exp(\kappa(x-1))$ , the inverse is  $\psi^{-1}(y) = 1 + \log(y)/\kappa$ .
# We want  $S(0) = \prod \psi_i(\theta_i) = 1 - f$ , so set  $\theta_i(0) = 1 + \log((1-f)^{1/n}) / \kappa_i$ .
function poisson_theta( $\kappa$ , f, n)
    1.0 + log((1 - f)^(1 / n)) /  $\kappa$ 
end
const  $\kappa_{hh}$ ,  $\kappa_{cm}$  = 3.0, 8.0
 $\theta_{hh}$  = poisson_theta( $\kappa_{hh}$ , seed_fraction_val, 2)
 $\theta_{cm}$  = poisson_theta( $\kappa_{cm}$ , seed_fraction_val, 2)

(sys_v, u0_v_default, ts_v, p_v) = build_multiplex_sir(layers_2; tspan=(0.0, 40.0))
u0_v = Dict{u0_v_default}
for var in collect(keys(u0_v))
    s = string(var)
    if occursin("theta_household", s)
        u0_v[var] =  $\theta_{hh}$ 
    elseif occursin("theta_community", s)
        u0_v[var] =  $\theta_{cm}$ 
    end
end
sol_v = solve(ODEProblem(sys_v, merge(u0_v, p_v), ts_v), Tsit5(); saveat = 0.5)
 $\theta_v$  = extract_theta(sol_v, sys_v, layer_names_2)
S_v = compute_S( $\theta_v$ , mean_degrees_2)
R_v = extract_R(sol_v, sys_v)
```

```

I_v = 1.0 .- S_v .- R_v

# --- Gillespie SSA on MultiplexNetwork -----
no_model_mp = OutbreakModel(
    [:S, :I, :R], [false, true, false],
    [OutbreakTransition(:S, :I, 1.0, :infection; via=:I),
     OutbreakTransition(:I, :R, 0.25, :spontaneous)];
    name = :SIR_multiplex)

builder = multiplex_graph_builder(N_sim, [
    (0.75, poisson_layer(3.0)), # household weight =  $\beta_{hh}$ 
    (0.25, poisson_layer(8.0)), # community weight =  $\beta_{cm}$ 
])

n_graphs_val = 5
nsims_per_g = 20
nsamples = n_graphs_val * nsims_per_g
S_runs = Matrix{Float64}(undef, nsamples, length(tgrid_val))
I_runs = similar(S_runs); R_runs = similar(S_runs)
rng = StableRNG(20240501)
row = 1
for gi in 1:n_graphs_val
    net = builder(rng)
    spec = OutbreakSpec(model = no_model_mp, network = net,
        initial = SeedFraction(:I  $\Rightarrow$  seed_fraction_val),
        tspan = (0.0, 40.0))
    ens = simulate_ensemble(spec; nsims = nsims_per_g,
        seed = 20240501 + 1000 * gi,
        algorithm = DirectSSA(),
        parallel = true)
    for traj in ens.trajectories
        for (j, t) in enumerate(tgrid_val)
            st = state_at(traj, t)
            S_runs[row, j] = st[no_model_mp.index_of[:S]] / N_sim
            I_runs[row, j] = st[no_model_mp.index_of[:I]] / N_sim
            R_runs[row, j] = st[no_model_mp.index_of[:R]] / N_sim
        end
        row += 1
    end
end
end
S_mean = vec(mean(S_runs; dims=1)); S_std = vec(std(S_runs; dims=1))
I_mean = vec(mean(I_runs; dims=1)); I_std = vec(std(I_runs; dims=1))

```

```
R_mean = vec(mean(R_runs; dims=1)); R_std = vec(std(R_runs; dims=1))
nothing
```

```
plt = plot(sol_v.t, S_v, color=:green, linewidth=2, label="S (EBCM)")
plot!(plt, sol_v.t, I_v, color=:red, linewidth=2, label="I (EBCM)")
plot!(plt, sol_v.t, R_v, color=:blue, linewidth=2, label="R (EBCM)")
plot!(plt, tgrid_val, S_mean, ribbon = S_std, color=:green,
      fillalpha=0.2, linealpha=0.6, linewidth=1, label="S (SSA)")
plot!(plt, tgrid_val, I_mean, ribbon = I_std, color=:red,
      fillalpha=0.2, linealpha=0.6, linewidth=1, label="I (SSA)")
plot!(plt, tgrid_val, R_mean, ribbon = R_std, color=:blue,
      fillalpha=0.2, linealpha=0.6, linewidth=1, label="R (SSA)")
xlabel!(plt, "Time"); ylabel!(plt, "Fraction of population")
title!(plt, "Two-layer multiplex SIR: EBCM vs Gillespie")
plt
```




Figure 10: Two-layer multiplex EBCM (lines, IC matched to a 1% seed) vs Gillespie SSA on a MultiplexNetwork host graph (mean $\pm 1\sigma$ ribbons across 5 graphs $\times$ 20 sims). Layers: Poisson(3) household with $\beta=0.75$ and Poisson(8) community with $\beta=0.25$; shared $\gamma=0.25$; $N=1000$.

With the IC matched, the deterministic EBCM trajectories track the stochastic ribbon means

closely across all three compartments, validating both the per-layer θ closure and the new `MultiplexNetwork` Direct SSA in `NetworkOutbreaks.jl`.