

Cross-Package SIS Comparison: EBCM vs Node-Based vs Stochastic

EdgeBasedModels.jl Vignette 16 — all-methods benchmark

Simon Frost

2026-05-13

Table of contents

Introduction	1
Setup	2
Stochastic ground truth (Gillespie SSA)	3
EBCM SIS	4
EBCM SIS with reinfection counting ($L = 1$)	4
Node-based moment closures	5
Standard pairwise: Keeling triple closure	5
Reinfection counting: $L = 1$	6
Motif closure	6
Neighbourhood closure: $n = 2$	6
Per-method comparison panels	7
Endpoint summary table	8
Transient diagnostic	9
Discussion	10
Reproducing this vignette	11

Introduction

This vignette places EdgeBasedModels.jl's EBCM alongside all four node-based moment closures from NodeBasedModels.jl and a Gillespie stochastic ground truth, on the same SIS benchmark as NBM vignette 13.

#	Method	Package
1	Gillespie SSA	NodeBasedModels
2	EBCM SIS (Miller)	EdgeBasedModels
3	EBCM edge-stratified reinfection ($L = 1$)	EdgeBasedModels
4	Standard pairwise (Keeling)	NodeBasedModels
5	Reinfection counting ($L = 1$)	NodeBasedModels
6	Motif closure ($m = 3$)	NodeBasedModels
7	Neighbourhood ($n = 2$)	NodeBasedModels

Setup

Note

$R_0 = 2$ anchor. For SIS on a k -regular configuration-model network, $R_0 = T(k - 1)$ with $T = \beta/(\beta + \gamma)$. We use $k = 4$, $\gamma = 0.25$, $T = 2/3$, hence $\beta = 0.5$, with 1% initial infection.

```
ENV["GKSwstype"] = "100" # offscreen GR output for reproducible Quarto renders
using EdgeBasedModels
import NodeBasedModels as NBM
using Graphs
using OrdinaryDiffEq
using Plots
using Random
using StableRNGs
using Statistics
using Printf
using Markdown

# Benchmark parameters – matched by the  $R_0=2$  invariant
N          = 500
k          = 4
γ_val      = 0.25
R0_target  = 2.0
κ_excess   = k - 1
T_val      = R0_target / κ_excess
β_val      = T_val * γ_val / (1 - T_val)
ε_val      = 0.01
tmax       = 120.0
ensemble   = 48
```

```

save_dt    = 1.0

tgrid      = collect(0.0:save_dt:tmax)
initial_infected = collect(1:round(Int, ε_val * N))

# Degenerate PGF for a k-regular network:  $\psi(z) = z^k$ 
regular_pgf(n::Int) = polynomial_pgf(vcat(zeros(n), [1.0]))

# Under the explicit-p convention, `seed_fraction` directly sets the
# initial node prevalence:  $\rho = \varepsilon_{\text{val}}$ , with  $\theta(0) = 1$ ,  $S(0) = (1-\rho) \cdot \psi(1)$ .
node_seed = ε_val

```

0.01

Stochastic ground truth (Gillespie SSA)

Host graph and ensemble seeds are identical to NBM vignette 13.

```

rng_host = StableRNG(20240301)
g         = random_regular_graph(N, k; rng = rng_host)
net       = NBM.GraphNetwork(g)

n_triangles = sum(triangles(g)) ÷ 3
n_p3_count  = sum((length(neighbors(g, v)) *
                    (length(neighbors(g, v)) - 1)) ÷ 2
                  for v in 1:nv(g)) - 3 * n_triangles

@printf("Host graph: N=%d, k=%d, triangles=%d, P3=%d\n",
        N, k, n_triangles, n_p3_count)

```

Host graph: N=500, k=4, triangles=3, P₃=2991

```

run_prev = zeros(ensemble, length(tgrid))
for r in 1:ensemble
    res = NBM.gillespie_sis(net;
                            infection_rate = β_val,
                            recovery_rate  = γ_val,
                            initial_infected = initial_infected,
                            tmax           = tmax,
                            seed           = 20240301 + r)

```

```

    for (i, t) in enumerate(tgrid)
        run_prev[r, i] = count(res(t)) / N
    end
end

gill_prev = vec(mean(run_prev; dims = 1))
gill_sd   = vec(std(run_prev; dims = 1))
@printf("Gillespie mean at t=%.1f: %.5f ± %.5f (1σ, n=%d)\n",
        tmax, gill_prev[end], gill_sd[end], ensemble)

```

Gillespie mean at t=120.0: 0.86725 ± 0.01739 (1σ, n=48)

EBCM SIS

The ODE for SIS on a k -regular network ($\psi(z) = z^k$) is $\dot{\theta} = -\beta\phi_I + \gamma(1 - \theta)$, $S = \psi(\theta)$, $I = 1 - S$, where $\phi_I = \theta - \psi'(\theta)/\psi'(1)$.

```

ebcm_sis = build_sis(regular_pgf(k), β_val, γ_val)

ic_sis = default_initial_conditions(ebcm_sis; seed_fraction = node_seed)
sol_sis = solve_epidemic(ebcm_sis;
                        tspan = (0.0, tmax),
                        init   = ic_sis,
                        saveat = save_dt,
                        reltol = 1e-8,
                        abstol = 1e-10)

I_sis = compartment(ebcm_sis, sol_sis, :I)
@printf("EBCM SIS endpoint I/N = %.5f\n", I_sis[end])

```

EBCM SIS endpoint I/N = 0.98232

EBCM SIS with reinfection counting ($L = 1$)

`build_sis_reinfection` tracks infection-count node strata $S_0, \dots, S_L$ and $I_1, \dots, I_L$ and separate edge densities such as `edge_S_0_I_1` and `edge_S_1_I_1`. Infection of each susceptible stratum is driven by its incident history-stratified infectious edges, so the aggregate $I(t) = \sum_p I_p(t)$ is no longer constrained to overlay the scalar EBCM trajectory.

```

L_reinf = 1
ebcm_reinf = build_sis_reinfection(regular_pgf(k),  $\beta$ _val,  $\gamma$ _val, L_reinf)
ic_reinf = default_initial_conditions(ebcm_reinf; seed_fraction = node_seed)
sol_reinf = solve_epidemic(ebcm_reinf;
                           tspan = (0.0, tmax),
                           init = ic_reinf,
                           saveat = save_dt,
                           reltol = 1e-8,
                           abstol = 1e-10)

I_reinf = compartment(ebcm_reinf, sol_reinf, :I)
@printf("EBCM + reinfection L=%d endpoint I/N = %.5f\n",
        L_reinf, I_reinf[end])

```

EBCM + reinfection L=1 endpoint I/N = 0.86957

Node-based moment closures

These calls mirror NBM vignette 13 exactly.

Standard pairwise: Keeling triple closure

```

sis = NBM.sis_model( $\tau$  = : $\beta$ )
hom = NBM.regular_network(k)

psys_pair = NBM.generate_pairwise(sis, hom, NBM.KeelingClosure();
                                   tspan = (0.0, tmax),
                                   seed_fraction =  $\epsilon$ _val)
sol_pair = NBM.solve_pairwise(psys_pair, Dict{: $\beta$  =>  $\beta$ _val, : $\gamma$  =>  $\gamma$ _val};
                               saveat = save_dt)
I_pair = sol_pair[psys_pair.singles[:I]]
@printf("Standard pairwise endpoint I/N = %.5f\n", I_pair[end])

```

Standard pairwise endpoint I/N = 0.86957

Reinfection counting: $L = 1$

As in the NBM companion vignette, the lifted system uses the same Keeling closure as the standard pairwise run. Since the homogeneous regular-network surrogate has $\phi = 0$, this is algebraically the Bernoulli/ordinary pair closure while keeping the comparison closure-matched.

```
psys_re = NBM.generate_pairwise(NBM.with_reinfection_counting(sis, 1),
                                hom, NBM.KeelingClosure();
                                tspan      = (0.0, tmax),
                                seed_fraction = ε_val)
sol_re   = NBM.solve_pairwise(psys_re, Dict{:β ⇒ β_val, :γ ⇒ γ_val};
                                saveat = save_dt)
I_re     = NBM.reinfection_totals(psys_re, sol_re)[:I]
@printf("NBM reinfection L=1 endpoint I/N = %.5f\n", I_re[end])
```

NBM reinfection L=1 endpoint I/N = 0.86957

Motif closure

The current motif implementation supports $k = 3$ but not the $k = 4$ regular network required for the feasible $R_0 = 2$ configuration-model anchor, so this panel is omitted from the refactored benchmark.

```
I_motif = nothing
```

Neighbourhood closure: $n = 2$

```
sys_nbr = NBM.generate_neighbourhood(sis, k, 2;
                                      β      = β_val, γ = γ_val,
                                      N      = 1.0,  ε = ε_val,
                                      tspan = (0.0, tmax))
sol_nbr  = NBM.solve_neighbourhood(sys_nbr; saveat = save_dt)
I_nbr    = NBM.neighbourhood_compartment(sys_nbr, sol_nbr, :I)
@printf("Neighbourhood n=2 endpoint I/N = %.5f\n", I_nbr[end])
```

Neighbourhood n=2 endpoint I/N = 0.86945

Verify all deterministic outputs land on the same time grid:

```
@assert length(I_sis)    = length(tgrid)
@assert length(I_reinf)  = length(tgrid)
@assert collect(sol_pair.t) = tgrid
@assert collect(sol_re.t)  = tgrid
@assert collect(sol_nbr.t) = tgrid
```


Per-method comparison panels

```
function comparison_panel(y, title, color, linestyle)
    p = plot(tgrid, gill_prev;
             ribbon = gill_sd,
             fillalpha = 0.16,
             lw = 2.4,
             color = :black,
             label = "Gillespie mean  $\pm 1\sigma$ ",
             xlabel = "Time",
             ylabel = "Prevalence I/N",
             title = title,
             legend = :bottomright,
             ylims = (0.0, 0.85))
    plot!(p, tgrid, y; lw = 2.4, color = color, ls = linestyle,
          label = title)
    return p
end

panels = [
    comparison_panel(I_sis, "EBCM SIS", :blue, :solid),
    comparison_panel(I_reinf, "EBCM reinfection L=1", :steelblue, :dash),
    comparison_panel(I_pair, "standard pairwise", :red, :dash),
    comparison_panel(I_re, "NBM reinfection L=1", :orange, :dashdot),
    comparison_panel(I_nbr, "neighbourhood n=2", :purple, :dash),
]

plot(panels...; layout = (3, 2), size = (1100, 900),
     plot_title = "Each deterministic approximation over Gillespie")
```

Each deterministic approximation over Gillespie



Endpoint summary table

```
ref = gill_prev[end]
rows = [
    ("Gillespie mean",          ref,          0.0),
    ("EBCM SIS (Miller)",      I_sis[end],   I_sis[end] - ref),
    ("EBCM + reinfection L=1", I_reinf[end], I_reinf[end] - ref),
    ("Standard pairwise (Keeling)", I_pair[end], I_pair[end] - ref),
    ("NBM reinfection L=1",    I_re[end],    I_re[end] - ref),
    ("Neighbourhood n=2",      I_nbr[end],   I_nbr[end] - ref),
]
lines = String[
    "| Method | Endpoint I/N | Signed deviation |",
```

```

    "|---|---:|---:|",
]
for (name, ep, dev) in rows
  push!(lines, @sprintf("| %s | %.5f | %+.5f |", name, ep, dev))
end
display(Markdown.parse(join(lines, "\n")))

```

Method	Endpoint I/N	Signed deviation
Gillespie mean	0.86725	+0.00000
EBCM SIS (Miller)	0.98232	+0.11507
EBCM + reinfection L=1	0.86957	+0.00232
Standard pairwise (Keeling)	0.86957	+0.00232
NBM reinfection L=1	0.86957	+0.00232
Neighbourhood n=2	0.86945	+0.00220

The Gillespie 1σ band at $t_{\max}$ is 0.01739. All deterministic methods that fall within $\pm 1\sigma$ of the stochastic mean can be considered well-calibrated on this benchmark.

Transient diagnostic

Both reinfection-counting constructions now alter transient aggregate prevalence. The EBM version does so through edge densities such as `edge_S_0_I_1`, while the NBM version does so through the full pairwise moment system. The max-over-time differences show where those corrections matter; endpoints can still be much closer after the $p = L$ bucket saturates.

```

function max_abs_delta(a, b)
  δ = abs.(a .- b)
  i = argmax(δ)
  return (value = δ[i], time = tgrid[i])
end

transient_rows = [
  ("EBCM SIS vs EBCM edge reinfection L=1",
   max_abs_delta(I_sis, I_reinf)),
  ("NBM pairwise vs NBM reinfection L=1",
   max_abs_delta(I_pair, I_re)),
  ("EBCM SIS vs Gillespie mean",
   max_abs_delta(I_sis, gill_prev)),
  ("NBM reinfection L=1 vs Gillespie mean",

```

```

    max_abs_delta(I_re, gill_prev)),
]

lines = String[
    "| Comparison | Max absolute difference | Time of max |",
    "|---|---:|---:|",
]
for (name, d) in transient_rows
    push!(lines, @sprintf("| %s | %.5f | %.1f |", name, d.value, d.time))
end
display(Markdown.parse(join(lines, "\n")))

```

Comparison	Max absolute difference	Time of max
EBCM SIS vs EBCM edge reinfection L=1	0.11275	97.0
NBM pairwise vs NBM reinfection L=1	0.04773	5.0
EBCM SIS vs Gillespie mean	0.11832	69.0
NBM reinfection L=1 vs Gillespie mean	0.01460	6.0

Discussion

EBCM accuracy. The SIR EBCM is exact on configuration-model networks in the $N \rightarrow \infty$ limit, but the scalar SIS builder used here is the standard one-dimensional EBCM-style closure for a non-monotone reinfection process. The edge-stratified reinfection model is a richer closure: it keeps the PGF-derived configuration-model triple factor but tracks history-stratified edge composition explicitly.

Why the comparison is non-trivial. On a finite random k -regular graph ($N = 500$) neither EBCM nor the node-based closures are exact. Repeated SIS infection and recovery induce local history correlations along edges; finite cycles and stochastic fluctuations add further deviations from the scalar edge-independence picture. Pairwise closures suffer from triple-approximation error, while motif and neighbourhood methods capture higher-order structure at the cost of larger state spaces. The endpoint table is therefore only one view; transient discrepancies are often more informative for this SIS case.

EBCM vs node-based. The EBCM is analytically tractable and cheap (one ODE for SIS). Node-based closures track pair and motif populations explicitly, making them more flexible for heterogeneous or clustered networks at the cost of a closure artefact. For the regular locally-tree-like host here the EBCM is the natural baseline; node-based methods become preferable when network-level statistics are of direct scientific interest.

Reproducing this vignette

- Host graph seed: `StableRNG(20240301)`.
- Gillespie ensemble: 48 runs, seeds 20240302–20240349.
- Initial infected: `collect(1:25)` (5 % of $N = 500$).
- ODE tolerances: `reltol = 1e-8`, `abstol = 1e-10` for the EBM solves and the package defaults for `solve_pairwise`, `solve_motif`, and `solve_neighbourhood`.
- Parameters identical to `NodeBasedModels.jl` vignette 13.

```
@printf("Julia %s\n", string(VERSION))
@printf("EdgeBasedModels path: %s\n", pathof(EdgeBasedModels))
@printf("NodeBasedModels path: %s\n", pathof(NBM))
```

Julia 1.12.5

EdgeBasedModels path: /Users/sdwfrost/Projects/edgebasedmodels/EdgeBasedModels.jl/src/EdgeBasedModels.jl

NodeBasedModels path: /Users/sdwfrost/Projects/edgebasedmodels/NodeBasedModels.jl/src/NodeBasedModels.jl